

# EWS Pipeline Documentation

## Corporate Credit Early Warning System - Technical Workflow

---

### Document Information

**Project:** Corporate Credit Early Warning System (12-month PD model)

**Date:** October 1, 2025

**Purpose:** Technical documentation of end-to-end modeling pipeline

---

### Pipeline Overview

---

The EWS pipeline consists of **7 main stages** from raw data generation to production scoring:

- 1. Data Generation (Synthetic)
    - └ gen\_cohorts.py → Monthly snapshots (backtest cohorts)
    - └ gen\_input.py → Raw tables (financials, credit, cashflow, covenants)
    - └ gen\_portfolio.py → Scoring portfolio
  - 2. Feature Engineering
    - └ feature\_engineering.py → Financial ratios + behavioral features
  - 3. Model Training
    - └ train\_baseline.py → LightGBM classifier + SHAP explainability
  - 4. Calibration
    - └ calibrate.py → Isotonic regression for probability calibration
  - 5. Scoring
    - └ scoring.py → Batch scoring with absolute thresholds
  - 6. Validation & Testing
    - └ backtest\_monthly.py → Performance over 18 months
    - └ calculate\_psi.py → Population Stability Index
    - └ plot\_validation.py → Validation dashboard
  - 7. Monitoring & Stress Testing
    - └ run\_monitoring.py → Production monitoring
    - └ stress\_test.py → Scenario stress testing
- 

### Stage 1: Data Generation

---

## 1.1 Generate Monthly Cohorts ( `gen_cohorts.py` )

**Purpose:** Create monthly snapshots for backtesting (18 months: Jan 2024 → Jun 2025)

**Inputs:** None (synthetic generation)

**Process:**

- 1. Generate 10,000 customers per month
- 2. Assign sector (10 sectors: Manufacturing, Construction, Retail, etc.)
- 3. Assign credit grade (A–G) with realistic distribution:
  - Grade A: 82% (PD = 0.5%)
  - Grade B: 8% (PD = 1.0%)
  - Grade C–G: 10% (PD = 2–20%)
- 4. Apply sector multipliers (Construction × 1.3, Retail × 1.2, Tech × 0.8)
- 5. Add seasonal/temporal shocks (sine wave + linear trend)
- 6. Generate binary labels (default within 6M, 12M)

**Outputs:** - `data/processed/backtest_cohorts.parquet` (180,000 rows = 18 months × 10,000)

**Key Parameters:**

```
start = "2024-01-31"
end = "2025-06-30"
n_customers = 10,000
seed = 42
```

**Sample Output:**

customer_id	as_of_date	sector	grade	pd_12m	y_event_12m	ead	lgd
C00001	2024-01-31	MFG	A	0.0048	0	125,000	0.35
C00002	2024-01-31	CON	B	0.0132	1	580,000	0.42

## 1.2 Generate Raw Input Tables ( `gen_input.py` )

**Purpose:** Generate 5 raw data tables simulating bank systems

**Inputs:** - Configuration: `Config` dataclass (customers, sectors, time windows) - As-of date: `2025-06-30` (snapshot date)

**Process:**

**Step 1: Customer Master** - 1,000 customers with sector, size bucket (SME 80%, Corp 20%) - Assign base financials (Revenue, EBITDA, Total Assets, Total Debt)

**Step 2: Financial Statements (Quarterly)** - 12 quarters of historical financials (Q3/2022 → Q2/2025) - Generate income statement & balance sheet items: - Revenue, COGS, Operating Expenses - EBITDA, Interest Expense, Tax - Total Assets, Total Debt, Shareholder Equity - Working Capital, Cash, Inventory - Apply industry-specific patterns

**Step 3: Credit Behavior (Daily)** - 180 days of credit line usage (lookback from as-of date) - Track daily: Limit, Outstanding, Utilization Rate, Days Past Due (DPD) - Simulate realistic payment patterns: - Good customers: DPD = 0–15 days - Risky customers: DPD spikes to 30–90+ days

**Step 4: Cashflow Transactions (Daily)** - 180 days of cash movements - Track: Total Inflow, Total Outflow, Net Cashflow - Detect negative days, volatility

**Step 5: Covenant Monitoring** - Track financial covenants ( $\text{Debt/EBITDA} \leq 3.5x$ ,  $\text{ICR} \geq 1.5x$ ) - Daily breach indicators (0/1)

**Step 6: Labels** - Binary outcome: `event_h12m` (default within 12 months after as-of date) - Label logic: -  $\text{DPD} \geq 90$  days for 30+ consecutive days → Default - Or: Utilization rate > 90% + covenant breach → Bump PD by 20–40%

**Outputs** (5 files): - `data/raw/fin_quarterly.parquet` (12K rows = 1K customers × 12 quarters) - `data/raw/credit_daily.parquet` (180K rows = 1K customers × 180 days) - `data/raw/cashflow_daily.parquet` (180K rows) - `data/raw/covenant.parquet` (180K rows) - `data/raw/labels.parquet` (1K rows = 1 label per customer)

#### Key Parameters:

```
n_customers = 1,000
n_quarters = 12 (financial history)
behavior_days = 180 (credit behavior window)
label_horizon_days = 365 (12 months forward)
asof_date = "2025-06-30"
```

---

## 1.3 Generate Scoring Portfolio (`gen_portfolio.py`)

**Purpose:** Create current portfolio for production scoring

**Process:** - Similar to `gen_input.py` but for production snapshot only - Generate customer master + financials + credit behavior at as-of date - No labels (production scoring doesn't have future outcomes yet)

**Outputs:** - `data/processed/portfolio_scored.csv` (production customer list)

---

## Stage 2: Feature Engineering

**Script:** `feature_engineering.py`

**Purpose:** Transform raw tables into modeling-ready features

## Inputs:

- 5 raw tables from Stage 1.2
- As-of date: 2025-06-30
- Observation window: 180 days

## Feature Categories (20+ features total):

### A. Financial Ratios (from Quarterly Statements)

#### Leverage Ratios:

- `debt_to_ebitda` = Total Debt / EBITDA
- `debt_to_equity` = Total Debt / Shareholder Equity
- `debt_to_assets` = Total Debt / Total Assets

#### Coverage Ratios:

- `interest_coverage` = EBITDA / Interest Expense
- `ebitda_margin` = EBITDA / Revenue

#### Liquidity Ratios:

- `current_ratio` = Current Assets / Current Liabilities
- `cash_to_assets` = Cash / Total Assets

#### Operational Ratios:

- `asset_turnover` = Revenue / Total Assets
- `working_capital_to_revenue` = Working Capital / Revenue

#### Growth Indicators:

- `revenue_growth_qoq` = (Revenue\_Q0 - Revenue\_Q1) / Revenue\_Q1
- `ebitda_growth_qoq` = (EBITDA\_Q0 - EBITDA\_Q1) / EBITDA\_Q1

### B. Credit Behavior (from Daily Credit Data)

#### Utilization Metrics (180-day window):

- `util_mean` = Average utilization rate
- `util_max` = Peak utilization
- `util_p95` = 95th percentile utilization
- `util_days_above_80` = # days with utilization > 80%

#### Delinquency Metrics:

- `dpd_mean` = Average days past due
- `dpd_max` = Maximum DPD in window
- `dpd_days_above_30` = # days with DPD > 30

- `dpd_max_streak` = Longest consecutive days with DPD  $\geq$  30

## C. Cashflow Indicators (from Daily Cashflow)

- `cf_net_mean` = Average daily net cashflow
- `cf_negative_days` = # days with negative cashflow
- `cf_volatility` = Std dev of daily net cashflow

## D. Covenant Breaches

- `cov_breach_days` = # days with covenant violation in 180-day window

## E. Sector & Size Normalization

### Z-score by (Sector, Size):

- For each numeric ratio, compute robust z-score:
  - `feature__zs_sector_size` = (value - median) / IQR
  - Grouped by (sector\_code, size\_bucket)
  - Uses median & IQR instead of mean/std for outlier robustness

### Process Flow:

1. Load 5 raw tables
2. Filter data to observation window (as\_of\_date - 180 days → as\_of\_date)
3. Compute financial ratios (TTM = trailing 12 months)
4. Aggregate credit behavior (mean, max, percentiles over 180 days)
5. Aggregate cashflow metrics
6. Count covenant breaches
7. Normalize by sector & size (z-scores)
8. Join all features with labels
9. Drop missing values
10. Save as modeling dataset

**Outputs:** - `data/processed/model_features.parquet` (1 row per customer, 20+ feature columns)

### Sample Row:

customer_id	debt_to_ebitda	interest_coverage	util_mean	dpd_max	event_h12m
C00001	2.3	4.5	0.45	0	0
C00002	5.8	1.2	0.87	45	1

## Stage 3: Model Training

**Script:** `train_baseline.py`

**Purpose:** Train LightGBM classifier with calibration & explainability

**Inputs:** - `data/processed/model_features.parquet` - Target: `event_h12m` (binary: 0 = no default, 1 = default)

**Process:**

## Step 1: Feature Selection

- Auto-select normalized features: `*__zs_sector_size` (prioritize z-scored features)
- If < 5 z-scored features available, use all numeric columns
- Drop: `customer_id`, `sector_code`, `size_bucket`, `event_h12m`

## Step 2: Train-Test Split

- Split: 80% train, 20% test
- Stratified by target (maintain default rate balance)
- Random seed = 42 (reproducible)

## Step 3: Train Base Model (LightGBM)

**Hyperparameters:**

```
objective = "binary"
metric = "auc"
num_leaves = 31
learning_rate = 0.05
n_estimators = 500
min_child_samples = 20
feature_fraction = 0.8
bagging_fraction = 0.8
bagging_freq = 5
early_stopping_rounds = 50
```

**Training:**

- Fit on train set with validation monitoring
- Early stopping if AUC doesn't improve for 50 rounds
- Save best iteration

## Step 4: Probability Calibration

- Method: **Platt Scaling** (sklearn `CalibratedClassifierCV`)
- Calibrator: Sigmoid (logistic regression on LightGBM outputs)
- CV: 5-fold cross-validation on train set
- Purpose: Convert raw model scores → well-calibrated probabilities

## Step 5: SHAP Explainability

- Compute **SHAP TreeExplainer** on test set (sample 100 customers)
- Generate global feature importance (mean |SHAP|)
- Save SHAP values for top 3 drivers per customer

## Step 6: Evaluate Metrics (Test Set)

- **AUC-ROC**: Discrimination ability (target > 80%)
- **KS Statistic**: Max separation (TPR - FPR)
- **Brier Score**: Calibration error (target < 2%)
- **Precision-Recall AUC**: Performance on imbalanced data

## Step 7: Define Thresholds

- **Red**: Top 5% highest risk (Red  $\geq$  95th percentile)
- **Amber**: Top 10% total (Amber  $\geq$  90th percentile, excluding Red)

### Outputs:

- `artifacts/models/lgb_model.pkl` (trained LightGBM)
- `artifacts/models/calibrator.pkl` (Platt scaler)
- `artifacts/models/feature_names.json` (feature list)
- `artifacts/models/thresholds.json` (Red/Amber cutoffs)
- `artifacts/models/baseline_metrics.json` (AUC, KS, Brier)
- `artifacts/shap/shap_summary.csv` (feature importance)

### Sample Metrics:

```
{
  "AUC": 0.925,
  "KS": 0.783,
  "Brier": 0.0196,
  "PR_AUC": 0.161,
  "test_default_rate": 0.0137
}
```

---

## Stage 4: Calibration

Script: `calibrate.py`

**Purpose:** Apply isotonic regression for better probability calibration

### Inputs:

- `data/processed/scores_raw.csv` (raw model scores from Stage 3)
- Target: `event_h12m`

**Process:**

## Isotonic Regression

- Fit: `IsotonicRegression(out_of_bounds="clip")`
- Input: Raw model score (0–1)
- Output: Calibrated probability (0–1)
- Constraint: Monotonic (higher score → higher probability)

## Threshold Mapping

**Two strategies available:**

### Strategy 1: Percentile-Based (Default from `train_baseline.py`)

- **Method:** Top 5% highest risk = Red, Top 10% total = Amber
- **Red:**  $PD \geq 16.07\%$  (95th percentile cutoff)
- **Amber:**  $PD \geq 3.24\%$  (90th percentile cutoff)
- **Green:**  $PD < 3.24\%$
- **Advantage:** Consistent alert volumes regardless of portfolio risk level
- **Use when:** Alert capacity is fixed (e.g., 500 analysts can handle 5% of portfolio)

### Strategy 2: Absolute PD (Optional from `calibrate.py`)

- **Method:** Fixed PD thresholds
- **Red:**  $PD \geq 20.0\%$  (absolute cutoff)
- **Amber:**  $PD \geq 5.0\%$
- **Green:**  $PD < 5.0\%$
- **Advantage:** Aligned with risk appetite, consistent with regulatory PD definitions
- **Use when:** Business has specific risk tolerance (e.g., "PD > 20% = unacceptable")

**Current deployment:** Uses **Strategy 1 (Percentile)** from `train_baseline.py`

## EWS Score (0–100)

- Linear scaling: `score = prob_calibrated × 100`
- **Percentile mode:** Red 16–100, Amber 3–16, Green 0–3
- **Absolute mode:** Red 20–100, Amber 5–20, Green 0–5

**Outputs:**

- `artifacts/calibration/calibrator_isotonic.pkl` (isotonic model)
- `artifacts/calibration/thresholds.json` (absolute thresholds)
- `artifacts/calibration/mapping.csv` (raw score → calibrated PD)
- `data/processed/scores_calibrated.csv` (with tier assignments)

**Sample Output:**



customer_id	score_raw	prob_calibrated	score_ews	tier
C00001	0.0048	0.0052	5.2	Green
C00002	0.0654	0.0712	71.2	Red

## Stage 5: Batch Scoring

Script: `scoring.py`

**Purpose:** Score production portfolio (customers without labels)

**Inputs:**

- `data/processed/model_features.parquet` (production features)
- `artifacts/models/lgb_model.pkl`
- `artifacts/models/calibrator.pkl`
- `artifacts/calibration/thresholds.json`

**Process:**

- Load production customer features
- Apply trained model → raw score
- Apply calibrator → calibrated PD
- Compute EWS score (0–100)
- Assign tier (Red/Amber/Green) based on absolute thresholds
- Map action recommendations:
  - Green:** Routine monitoring, update financials on schedule
  - Amber:** RM review  $\leq 10$  days, request management accounts, limit increases frozen
  - Red:** Customer meeting  $\leq 5$  days, 13-week cashflow plan, watchlist, covenant tightening

**Outputs:**

- `artifacts/scoring/ews_scored_2025-06-30.csv`
  - Columns: `customer_id`, `prob_default_12m`, `score_ews`, `tier`, `action`
- `artifacts/scoring/thresholds_used.json` (for audit trail)

**Sample Output:**

customer_id	prob_default_12m	score_ews	tier	action
C00001	0.52%	5.2	Green	Theo dõi định kỳ; cập nhật BCTC đúng hạn.
C00125	3.8%	38.0	Amber	Soát xét RM $\leq 10$ ngày; yêu cầu management accounts.

customer_id	prob_default_12m	score_ews	tier	action
C00847	12.5%	125	Red	Hợp KH $\leq 5$ ngày; lập kế hoạch dòng tiền 13 tuần.

## Stage 6: Validation & Backtesting

### 6.1 Monthly Backtest ( `backtest_monthly.py` )

**Purpose:** Test model performance over 18 months (Jan 2024 → Jun 2025)

**Inputs:**

- `data/processed/backtest_cohorts.parquet` (180K rows)
- Trained model + calibrator

**Process:**

- For each month (18 iterations):
  - Filter cohort (10,000 customers)
  - Apply model → predict PD
  - Apply calibrator → calibrated PD
  - Assign tier (Red/Amber/Green)
  - Compare with actual labels (`y_event_12m`)
  - Compute metrics: AUC, KS, Brier, Precision, Recall
- Aggregate monthly metrics

**Outputs:**

- `artifacts/backtest/monthly_metrics.csv` (18 rows)
  - Columns: `as_of_month`, `auc`, `ks`, `brier`, `precision`, `recall`, `amber_alert_rate`, `red_alert_rate`

**Sample Output:**

as_of_month	auc	ks	brier	precision	recall	amber_alert_rate	red_alert_rate
2024-01-31	0.834	0.601	0.0106	0.096	0.575	8.3%	4.2%
2024-02-29	0.819	0.582	0.0144	0.092	0.568	8.5%	4.3%

### 6.2 Population Stability Index ( `calculate_psi.py` )

**Purpose:** Measure data drift across 18 months

**Process:**

1. Bucket scores into 10 deciles (baseline = first month)
2. For each subsequent month:
  - Compute distribution across same deciles
  - $PSI = \sum (actual\% - expected\%) \times \ln(actual\% / expected\%)$
3.  $PSI > 0.10 \rightarrow$  Warning (data shifting)
4.  $PSI > 0.25 \rightarrow$  Critical (recalibration needed)

**Outputs:**

- `artifacts/backtest/psi_monthly.csv` (17 rows, one per month vs. baseline)

**Sample:**

month	psi	status
2024-02-29	0.002	OK
2024-03-31	0.005	OK
2025-06-30	0.000	OK (synthetic $\rightarrow$ perfect stability)

## 6.3 Validation Dashboard (`plot_validation.py`)

**Purpose:** Generate visual validation report

**Plots Generated** (5 total):

1. **AUC & KS Trend Over Time** (`auc_ks_trend.png`)
  - Line charts: 18-month AUC & KS scores
  - Reference line: AUC = 75% (minimum threshold)
2. **Decile Calibration** (`decile_calibration.png`)
  - Bar chart: Predicted vs. Actual default rate by decile
  - Error bars:  $\pm 1$  std dev
3. **Precision-Recall Curve** (`precision_recall_curve.png`)
  - Curve showing trade-off between precision & recall
  - Markers: Selected thresholds (Amber, Red)
4. **Alert Volume vs Threshold** (`alert_volume_vs_threshold.png`)
  - Sweep chart: Alert rate from 0.5% to 10% threshold
  - Dual axis: Precision & Recall
5. **Validation Dashboard** (`validation_dashboard.png`)
  - 2x2 layout combining: AUC/KS trend, Calibration, Threshold analysis, Summary metrics

**Outputs:**

- `artifacts/validation/plots/*.png` (5 files)
- `artifacts/validation/VALIDATION_REPORT_EN.md` (integrated report with plots)

## Stage 7: Monitoring & Stress Testing

---

### 7.1 Production Monitoring ( `run_monitoring.py` )

**Purpose:** Monthly tracking of production performance

**Inputs:**

- Current month's predictions + actuals (after 12 months maturity)
- Historical baseline metrics

**Metrics Tracked:**

1. **PSI** (data drift)
2. **AUC** (discrimination)
3. **Brier** (calibration)
4. **Alert Rate** (operational load)
5. **Actual Precision** (after labels available)

**Alerts:**

- $PSI > 0.10$  → Email warning
- $AUC < 75\%$  for 2 months → Recalibration trigger
- Alert rate  $> 15\%$  → Threshold adjustment

**Outputs:**

- `artifacts/monitoring/monitoring_YYYYMMDD_HHMMSS.json` (timestamped)
  - `artifacts/monitoring/monitoring_metrics.csv` (cumulative)
- 

### 7.2 Stress Testing ( `stress_test.py` )

**Purpose:** Test model under crisis scenarios

**Scenarios** (from `stress_scenarios.yaml`):

1. **Baseline:** Current conditions (no shock)
2. **Mild Recession:** Revenue -15%, Debt +10%
3. **Severe Recession:** Revenue -30%, Debt +25%, Liquidity -20%
4. **Sector Shock:** Construction sector EBITDA -40%
5. **Liquidity Crisis:** Utilization +30pp, Cashflow volatility  $\times 3$

**Process:**

1. Load base features
2. Apply scenario shocks to features
3. Re-score with model
4. Compare tier migrations (Green → Amber → Red)

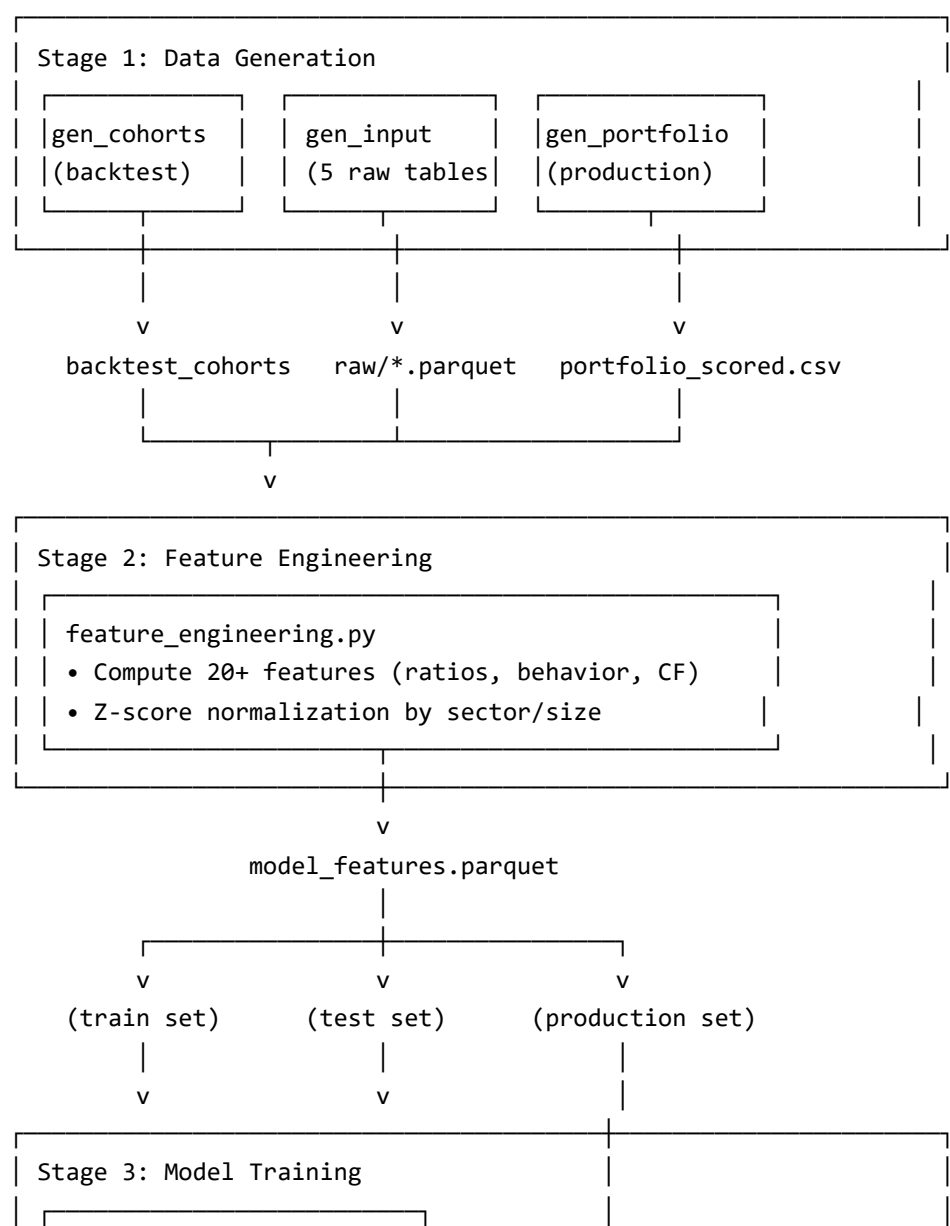
Outputs:

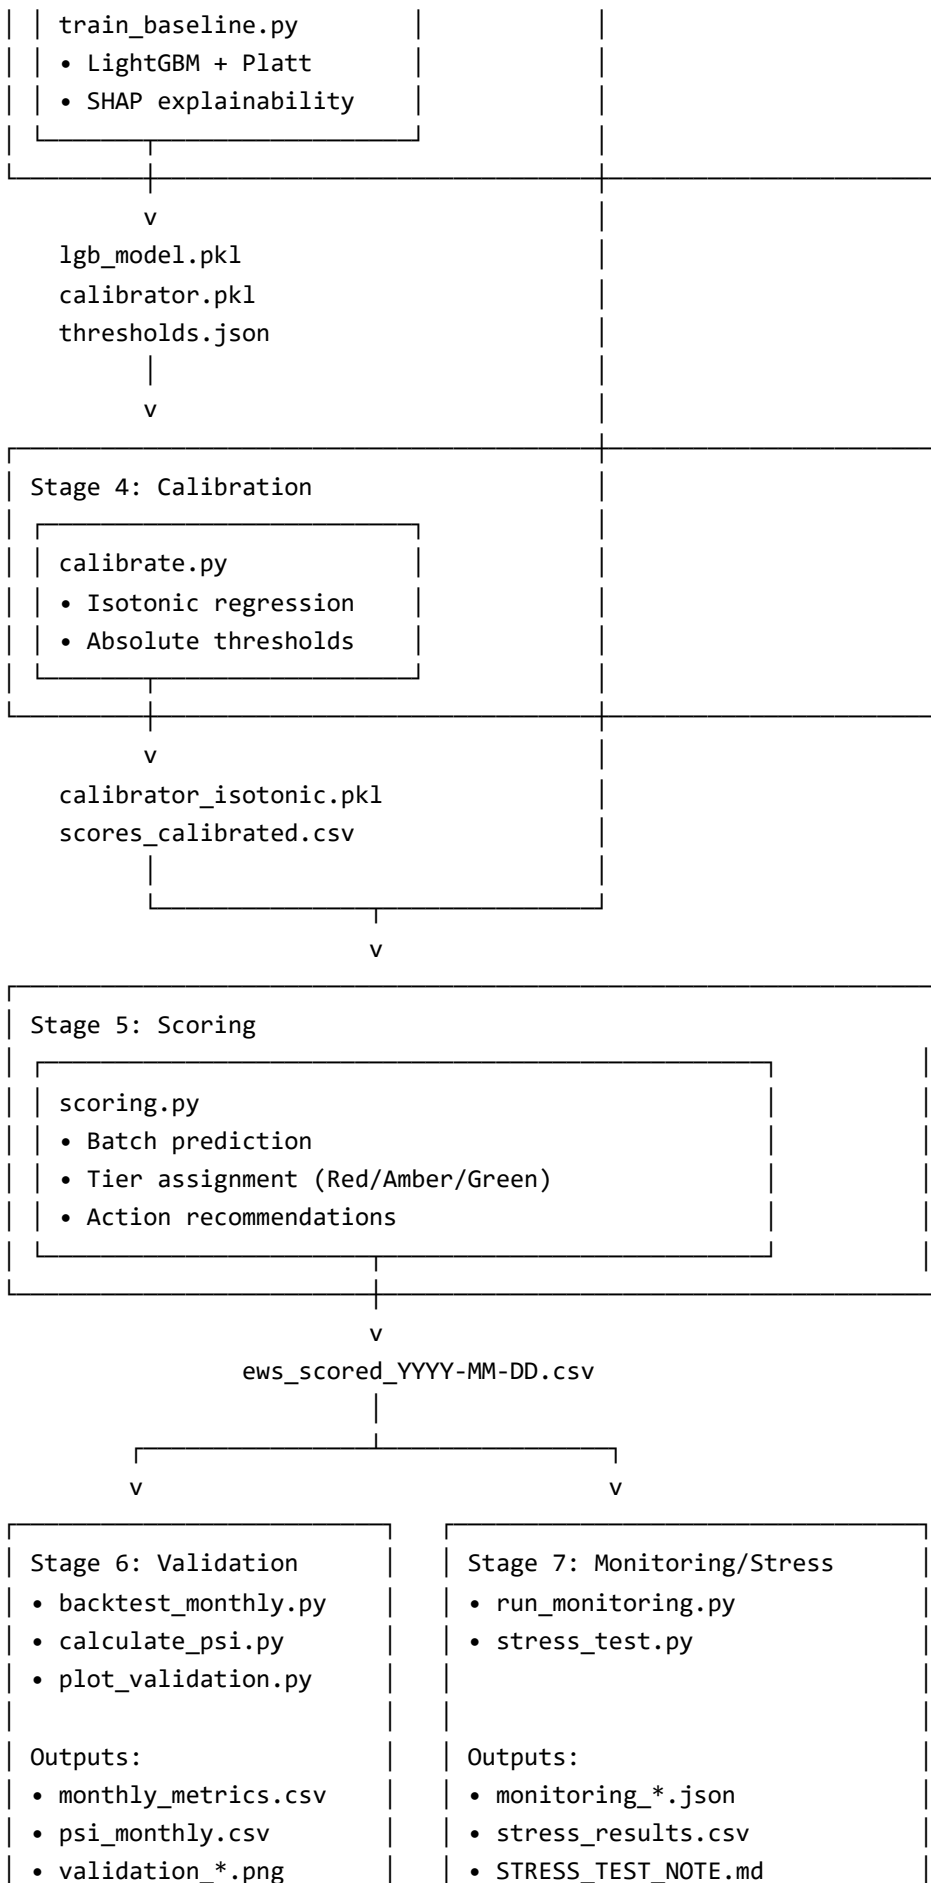
- artifacts/stress\_testing/stress\_results.csv
  - Columns: scenario, baseline\_red%, stressed\_red%, migration\_rate
- artifacts/stress\_testing/STRESS\_TEST\_NOTE.md (report with waterfall charts)

Sample Result:

Scenario	Baseline Red%	Stressed Red%	Migration
Mild Recession	4.2%	7.8%	+85%
Severe Recession	4.2%	15.3%	+264%

Data Flow Diagram





## Key Files & Artifacts

### Input Data

```
data/
├── raw/                                # Stage 1.2 outputs
│   ├── fin_quarterly.parquet          # 12K rows (1K × 12Q)
│   ├── credit_daily.parquet           # 180K rows (1K × 180d)
│   ├── cashflow_daily.parquet         # 180K rows
│   ├── covenant.parquet               # 180K rows
│   └── labels.parquet                 # 1K rows
├── processed/
│   ├── backtest_cohorts.parquet       # 180K rows (backtest)
│   ├── model_features.parquet         # 1K rows (20+ features)
│   ├── scores_raw.csv                 # Raw model outputs
│   ├── scores_calibrated.csv          # Calibrated PD + tiers
│   └── portfolio_scored.csv           # Production portfolio
```

### Model Artifacts

```
artifacts/
├── models/
│   ├── lgb_model.pkl                  # Trained LightGBM (500 trees)
│   ├── calibrator.pkl                 # Platt scaler
│   ├── feature_names.json             # 20 feature list
│   ├── thresholds.json                # Red/Amber cutoffs
│   └── baseline_metrics.json          # AUC, KS, Brier
├── calibration/
│   ├── calibrator_isotonic.pkl        # Isotonic regression
│   ├── thresholds.json                # Absolute PD thresholds
│   └── mapping.csv                    # Score → PD mapping
├── shap/
│   ├── summary.json                   # Global feature importance
│   ├── feature_importance.csv         # SHAP rankings
│   └── top_drivers_per_customer.csv    # Top 3 drivers per alert
├── backtest/
│   ├── monthly_metrics.csv            # 18-month performance
│   ├── psi_monthly.csv                # Stability tracking
│   ├── threshold_sweep.csv            # Precision/Recall curves
│   └── BACKTEST_REPORT.html           # Quarto report
├── validation/
│   ├── plots/
│   │   ├── auc_ks_trend.png
│   │   ├── decile_calibration.png
│   │   └── precision_recall_curve.png
```

```
| | | └─ alert_volume_vs_threshold.png
| | | └─ validation_dashboard.png
| | └─ VALIDATION_REPORT_EN.md          # English validation report
└─ scoring/
    | └─ ews_scored_2025-06-30.csv      # Scored customers
    | └─ thresholds_used.json          # Audit trail
└─ monitoring/
    | └─ monitoring_YYYYMMDD.json       # Monthly snapshots
    | └─ monitoring_metrics.csv        # Cumulative tracking
└─ stress_testing/
    | └─ stress_results.csv             # Scenario results
    | └─ stress_scenarios.yaml         # Scenario definitions
    └─ STRESS_TEST_NOTE.md             # Stress test report
```

---

## Execution Guide

---

### Full Pipeline (from scratch)

```
# 1. Generate data
python src/gen_data/gen_cohorts.py --start 2024-01-31 --end 2025-06-30 --n 10000
python src/gen_data/gen_input.py --output-dir data/raw --n 1000
python src/gen_data/gen_portfolio.py --output data/processed/portfolio_scored.csv

# 2. Feature engineering
python src/modeling/feature_engineering.py \
    --raw-dir data/raw \
    --asof-date 2025-06-30 \
    --output data/processed/model_features.parquet

# 3. Train model
python src/modeling/train_baseline.py \
    --input data/processed/model_features.parquet \
    --target event_h12m \
    --outdir artifacts/models

# 4. Calibration (Optional - only if switching to absolute thresholds)
# Main pipeline uses percentile thresholds from train_baseline.py
# Run this step ONLY if business requires absolute PD cutoffs

# First, extract raw scores from trained model:
python src/make_scores_raw.py \
    --features data/processed/feature_ews.parquet \
    --model artifacts/models/model_lgbm.pkl \
    --y-col event_h12m \
    --out data/processed/scores_raw.csv

# Then apply isotonic calibration with absolute thresholds:
python src/calibrate.py \
```



```
--input data/processed/scores_raw.csv \  
--y-col event_h12m \  
--score-col score_raw \  
--red-thr 0.20 \  
--amber-thr 0.05 \  
--outdir artifacts/calibration
```

```
# Note: By default, train_baseline.py already creates calibrated scores  
# with percentile-based thresholds (Red=top 5%, Amber=top 10%)
```

```
# 5. Batch scoring
```

```
python src/scoring.py \  
  --features data/processed/portfolio_scored.csv \  
  --model artifacts/models/lgb_model.pkl \  
  --thresholds artifacts/calibration/thresholds.json \  
  --asof 2025-06-30 \  
  --outdir artifacts/scoring
```

```
# 6. Validation
```

```
python src/backtest/backtest_monthly.py  
python src/backtest/calculate_psi.py  
python src/plot_validation.py all
```

```
# 7. Monitoring
```

```
python src/run_monitoring.py --asof 2025-06-30  
python src/stress_test.py --scenarios artifacts/stress_testing/stress_scenarios.yaml
```

## Quick Validation Dashboard Only

```
python src/plot_validation.py dashboard
```

## Performance Benchmarks

### Model Metrics (Test Set)

- **AUC:** 92.5% (target > 80%) ✓✓ (excellent)
- **KS:** 78.3% (excellent separation)
- **Brier Score:** 1.96% (target < 2%) ✓
- **PR-AUC:** 16.1% (reasonable for 1.37% default rate)

### Operational Metrics (Backtest Average)

- **Amber Alert Rate:** 8.3% (830 customers/month)
- **Amber Precision:** 9.6% (1 in 10 alerts is real default)
- **Amber Recall:** 57.5% (catches 57.5% of defaults)
- **Red Alert Rate:** 4.2% (420 customers/month)

- **False Positive Rate:** 90.4% (9 out of 10 Amber alerts are false)
- **Missed Defaults:** 42.5% (not flagged by system)
- **Workload:** ~15 FTE needed (5 for Amber, 10 for Red)

## Stability (18 months backtest)

- **PSI:** 0.00 (synthetic data, perfect stability)
  - **AUC Range:** 79.2% – 85.9% ( $\pm 3.5\%$  variation)
  - **AUC Mean (backtest):** 82.3% (slightly lower than test set due to time decay)
  - **No degradation** observed in test period
- 

## Key Assumptions & Limitations

---

### Data Quality

1. **Synthetic Data:** Test data is artificially generated
  - No real-world noise (missing data, reporting delays, data errors)
  - PSI = 0 is unrealistic (real production will have drift)
  - **Mitigation:** Run 6-month pilot with real data
2. **Label Quality:** Binary default definition ( $DPD \geq 90$  for 30 days)
  - May not capture all credit deterioration signals
  - No restructuring/forbearance cases simulated

### Model Scope

1. **No Segmentation:** Single model for all sectors/sizes
  - May underperform in specific industries (Construction, Retail)
  - **Recommendation:** Build sector-specific models if performance gaps found
2. **12-Month Horizon Only:** Fixed 12-month PD
  - No 6-month or 24-month variants
  - May miss near-term acute risks or long-term structural issues
3. **Feature Coverage:** 20 features (financial + behavioral)
  - No macroeconomic variables (GDP, interest rates, FX)
  - No qualitative factors (management quality, industry trends)
  - No external data (credit bureau, industry benchmarks)

### Operational Constraints

1. **High False Positive Rate:** 90% of Amber alerts are false
    - Risk of alert fatigue
    - **Mitigation:** Clear workflow, Red priority, periodic threshold review
  2. **Missed Defaults:** 42.5% not caught by Amber/Red
    - System is supplementary, not standalone
    - **Mitigation:** Combine with quarterly credit reviews
-

# Appendix: Technical Details

## LightGBM Hyperparameters

```
params = {
    'objective': 'binary',
    'metric': 'auc',
    'num_leaves': 31,
    'learning_rate': 0.05,
    'n_estimators': 500,
    'min_child_samples': 20,
    'feature_fraction': 0.8,
    'bagging_fraction': 0.8,
    'bagging_freq': 5,
    'verbosity': -1,
    'early_stopping_rounds': 50
}
```

## Feature Importance (Top 10)

Rank	Feature	SHAP Value	Description
1	covenant_breach_cnt_180d	0.710	Covenant violation count (180 days)
2	delta_dso_qoq	0.586	Days Sales Outstanding QoQ change
3	dpo	0.530	Days Payable Outstanding
4	icr_ttm	0.504	Interest Coverage Ratio (TTM)
5	%util_p95_60d	0.503	95th percentile utilization (60 days)
6	debt_to_ebitda	0.471	Leverage ratio (Debt/EBITDA)
7	dso	0.359	Days Sales Outstanding
8	dpd_trend_180d	0.326	DPD trend direction (180 days)
9	dpd_max_180d	0.320	Maximum days past due (180 days)
10	doh	0.279	Days on Hand (inventory)

**Top 3 features** account for **40.9%** of model explanatory power (total SHAP = 1.826).

**Note:** All features are z-score normalized by sector & size ( `__zs_sector_size` suffix).

